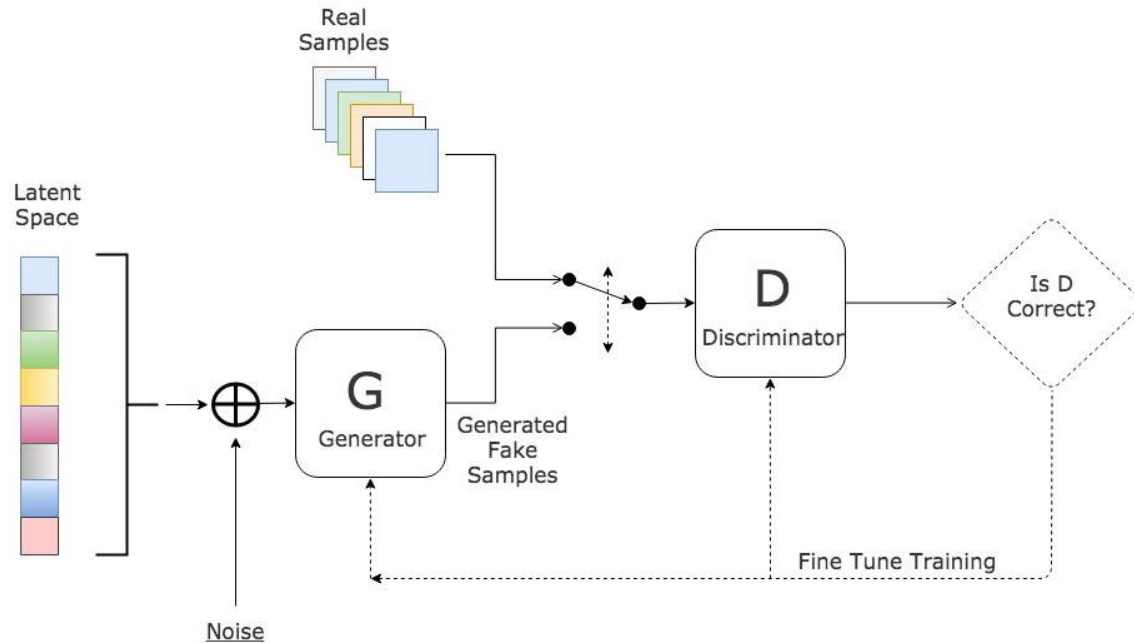


Generative Adversarial Networks (GANs)

Group 3


What is a GAN?

Generative Adversarial Network



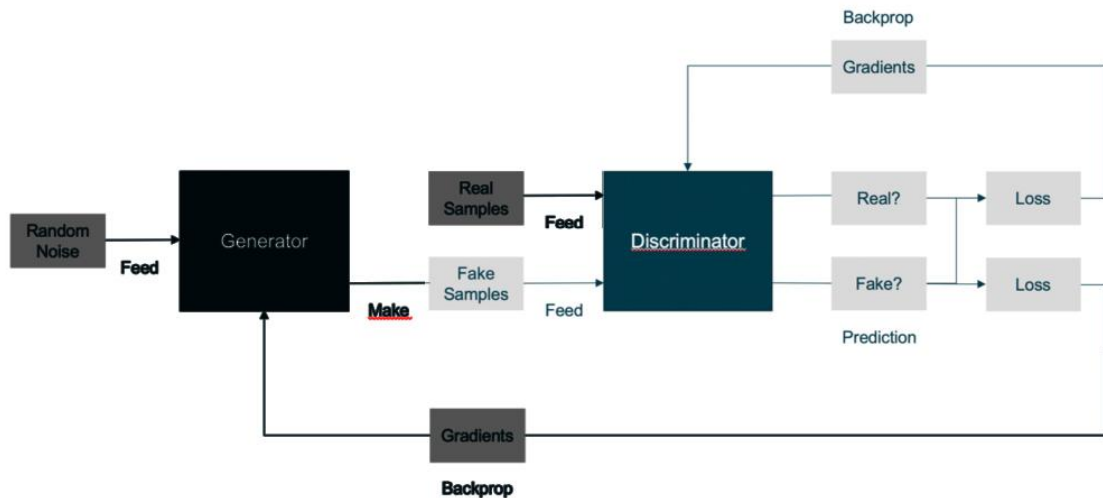
Activation Function

LeakyRelu.



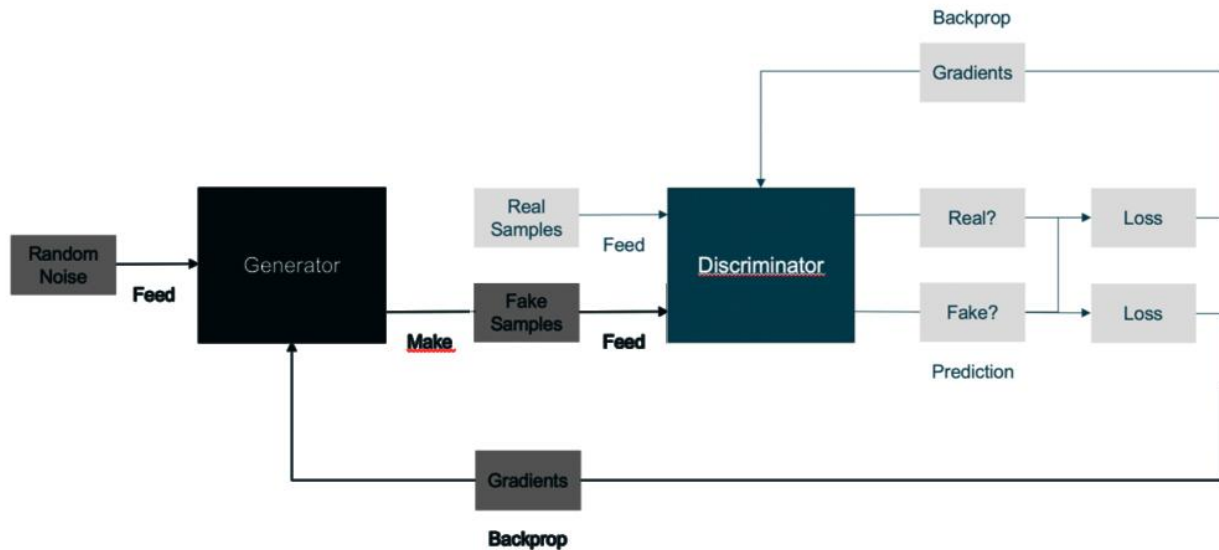
Training a GAN

Freeze the Generator and train the Discriminator once on generated samples. You get `d1_loss`



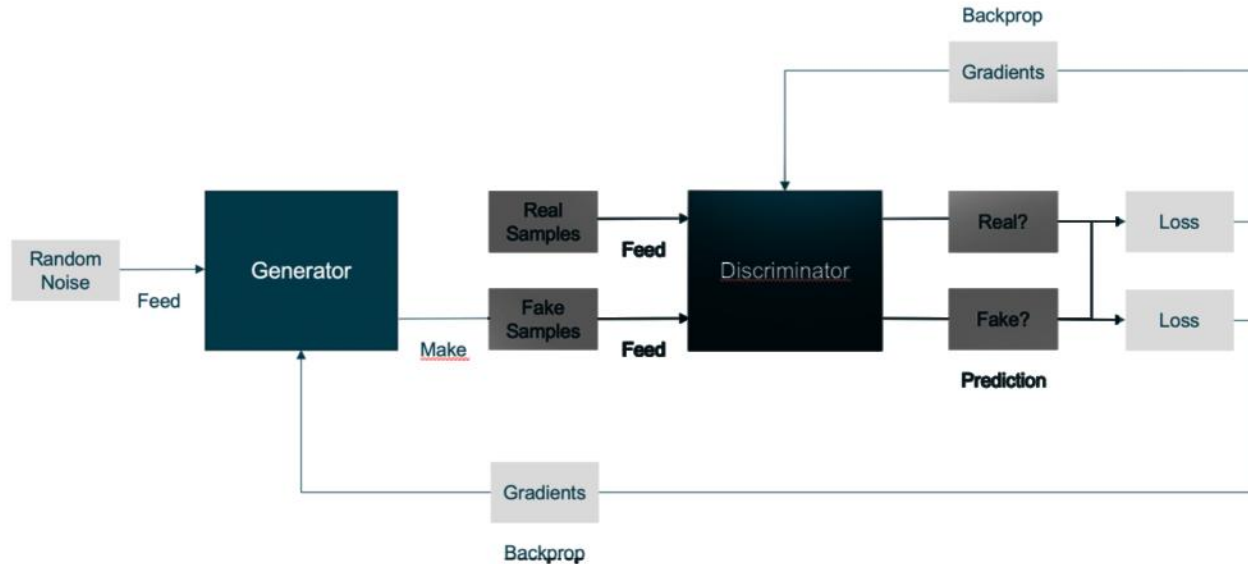
Training a GAN

Freeze the Generator and train the Discriminator once on real samples. You get `d2_loss`



Training a GAN

Freeze the Discriminator and train the Generator with Latent vector. You get g_loss



Training a GAN

The goal is to minimize generator loss and maximize discriminator loss.


Meaning:

If g loss is small, it means the generator is able to produce very realistic images

If d loss is high, it means it is not able to classify images into real or fake, meaning the generator is doing a good job.

Data Description

- This is a dataset consisting of 21551 anime faces
- All images are resized to 64 * 64 for the sake of convenience.



Baseline Model Description

- Generator

- Input layer
- Conv2DTranspose layer
- Batchnormalization layer
- LeakyReLU()
- Conv2D layer
- Tanh()

- Discriminator

- Input layer
- Conv2D layer
- LeakyReLU()
- Dropout(rate=0.3)
- Flatten layer
- Fully connected layer

Generated Samples of Baseline Model



Optimization Methods


- **Adjust the learning rate (0.0001, 0.0002, 0.0005)**
- **Add more convolutional layers**
- **Use a different loss function (MeanSquaredError, Hinge Loss)**
- **Use a different optimizer (SGD,RMSprop)**
- **Changing Kernel Size**
- **LeakyRelu Value(0.05,0.1,0.5,0.2)**
- **Adding more Filters**
- **Increase the number of training epochs**
- **BatchNormalization in discriminator**
- **Change activation function in generator**
- **Adding noise in discriminator**

Losses


Generator

Discriminator(Generated Samples)

Discriminator(Real Samples)



Results



LR:0.0001



LR:0.0002



LR:0.0005

Results




SGD



RMSprop


Results



Hinge loss



MeanSquaredError



kl_divergence_loss

Results

LeakyRelu Value



0.05



0.1



0.5




0.2

Results

Adding noise in discriminator




adding noise after first input layer



adding noise after each layer

Results

Filter, epochs



Optimized Model Description


- Generator

- Input layer
- Conv2DTranspose layer
- Batchnormalization layer
- ReLU()
- Conv2D layer
- Tanh()

- Discriminator

- Input layer
- Conv2D layer
- Batchnormalization layer
- LeakyReLU()
- Dropout(rate=0.3)
- Flatten layer
- Fully connected layer

Architecture




Optimized Model Result





```
df['d1_loss'].plot()
```

<Axes: >




```
df['d2_loss'].plot()
```

<Axes: >



```
df['g_loss'].plot()
```

<Axes: >



Best Result



Reference

<https://machinelearningmastery.com/how-to-develop-a-conditional-generative-adversarial-network-from-scratch/>

<https://github.com/hwalsuklee/tensorflow-generative-model-collections>

<https://machinelearningmastery.com/how-to-evaluate-generative-adversarial-networks/>

<https://www.hindawi.com/journals/misy/2022/9005552/>

<https://arxiv.org/pdf/1801.09195.pdf>

<https://github.com/nikhilroxtomar/DCGAN-on-Anime-Faces/blob/master/gan.py>

<https://arxiv.org/abs/1606.03498>

<https://pyimagesearch.com/2020/11/16/gans-with-keras-and-tensorflow/>

https://proceedings.neurips.cc/paper_files/paper/2018/file/90365351ccc7437a1309dc64e4db32a3-Paper.pdf

<https://paperswithcode.com/method/label-smoothing>

<https://arxiv.org/pdf/1511.06434v2.pdf>

<https://www.kaggle.com/datasets/soumikrakshit/anime-faces>